Research Review

Casey Battaglino

January 22, 2014

1 Introduction

This document presents one section for each topic I've been exploring. Each section is structured as follows:

1. **Overview**: A concise introduction to existing work on the topic, references, and motivation.

2. **Results / Implementation**: Any original experiments, tests, visualizations, and implementations related to the topic.

3. **Thoughts**: Any original musings or ideas on extending the existing work.

The main topics (sections) are as follows:

1. **Distributed-Memory (Re-)Streaming Graph Partitioning**: Exploring streaming partitioning, implementing re-streaming partitioning (and then discovering it's already been done), implementing a shared-memory and distributed-memory streaming partitioner, discussion of quality of partitions.

2. **Optimizing Distributed Scale-Free Graph Traversals**: Exploring a scheme for distributed-memory graph traversals that involves low-degree and high-degree nodes in different communication patterns (and leveraging the fact that low-degree nodes can be partitioned, in isolation, at much higher quality, for scale-free graphs).

3. **Formalizing Characteristics of Graph Traversals**: Leveraging statistical qualities of random graph models to quantify the performance of algorithms operating upon them.

4. **Other Open Questions**: Other thoughts and possible directions.

Finally, I reiterate over the main "TODO" points, which we can prioritize and add to as we see fit.

2 Distributed-Memory (Re-)Streaming Graph Partitioning

2.1 Overview

Streaming graph partitioning has only been introduced in the last couple of years [**?**, **?**, **?**], as graphs grow to the point where they may not fit into memory, or must be distributed to compute nodes on the fly. In the streaming model, input data (vertices) arrive sequentially from a generating source (such as a web-crawler), and must be partitioned as they arrive.

The order in which vertices are processed determines the final partition assignment. For instance, a web crawler might generate vertices in an order that corresponds to a Breadth-First or Depth-First traversal of the web, or we may receive vertices in a random order. An analysis of streaming algorithms may also consider an *adversarial* ordering that produces the worst possible results [**?**].

Assuming vertices arrive in some order, a heuristic makes a partition decision, given vertex v, and capacity constraint C (where C is generally $\approx \frac{(\epsilon + |V|)}{n}$ Stanton explores a number of heuristics (from simple to complex), and finds that a simple weighted greedy approach is most effective [**?**]. This heuristic is as follows:

Weighted Deterministic Greedy (WDG): Assign v to partition that it has most edges to, weighted by the relative size of each partition (weight function can be linear or exponential) More formally:

$$ind = \underset{i \in [k]}{\arg\max}\{|P_t(i) \cap \Gamma(v)|w(t,i)\}$$

$w(t,i) = 1$ for unweighted greedy

$w(t,i) = 1 - \frac{|P^t(i)|}{C}$ for linear weighted

$w(t,i) = 1 - \exp\{|P^t(i)| - C\}$ for exponentially weighted

This heuristic proves surprisingly effective for scale-free graphs, and can be competitive with leading offline implementations such as METIS, while operating at a much faster speed. The low diameter of scale-free graphs seems to make them easier to partition in this way (because as we process more and more nodes, the chance that an incoming node has neighbors that have already been partitioned becomes very high. This is not true for spatial, structured graphs).

Partitioning a 26GB Twitter follower graph can take nearly a day using METIS, but can take a matter of minutes using a streaming algorithm, with similar partition quality [**?**].

2.2 Results / Implementation

2.2.1 Shared-Memory Streaming Partitioner

I implemented the generalized WDG partitioner FENNEL, and tested it on a wide variety of real-world and synthetic graphs. Our initial 'test-bed' was implemented in MATLAB, because many important properties of graphs can be easily extracted using vector notation on their adjacency matrices. As we scaled to larger graphs, the outer for-loop in the code (which iterates over each vertex) began to take an infeasible amount of time, so I implemented a more fully-featured test-bed in C. The MATLAB and C code is located at `https://github.com/cjbattagl/graphpart/`. The

C code makes significant use of BeBoP's Sparse Matrix Converter: `http://bebop.cs.berkeley.edu/smc/`.

Our implementation can read graphs in three formats (Harwell-Boeing, Matlab, and Matrix-Market Formats). This allows us to read in the entire SNAP graph archive [**?**]. It can read in symmetric (undirected) and unsymmetric (directed) graphs. Once the graph is read in, we convert it from unordered triplet format to Compressed Sparse Row format, which allows us to iterate through the neighbor list of any node in linear time, as well as compute the degree of any node in constant time.

Once we have our graph stored as a matrix in CSR form, we can perform a simulation of streaming $k-$partitioning. First we generate the traversal order of the graph (a random vertex order, which we generate using a Knuth Shuffle). We maintain a data structure that contains the assignment of every assigned vertex to its partition. Thus, for a given vertex, we can compute the number of edges from that vertex to each partition by iterating over each edge incident to the vertex. We also maintain a running total of vertices in each partition. This process lets us compute the FENNEL algorithm in $O(|E|) = O(nnz(A))$ time.

We measured the overall quality of partitions by the *fraction of cut edges* λ.

$$\lambda = \frac{\text{Number of edges cut by partition}}{\text{Total number of edges}} \tag{1}$$

In general, we can compare this to the expected quality of a random $k-$partition. To estimate this we just assume that each block has the same density of nonzeros, so we compute the fraction of nonzeros expected to be in non-diagonal blocks:

$$\lambda_r = \frac{k^2 - k}{k^2} = \frac{k-1}{k} \tag{2}$$

Any partitioner that produces a partition with $\lambda < \lambda_r$ has thus improved the parallel locality of the graph ordering, so we present $\lambda_{r,k}$ along with our results.

In Figure **??** we show the main properties of graphs in the SNAP database, as well as the performance of our streaming partitioner on them, for $k = 2$ and $k = 8$.

2.2.2 Distributed Memory Streaming Partitioner

I implemented a streaming partitioner for distributed memory using MPI.

1. The vertices of the graph are logically mapped in a block partitioning to each process.

2. The vertices (row pointers) and associated edges (nnzs) are communicated to each process.

3. Partitioning begins. After a chunk of vertices has been partitioned, its partition is *scattered* to all vertices so that the greedy function can be computed.

4. Upon termination, partitioning can either restart (re-streaming), or the graph can be redistributed.

Figure 1: Basic properties of graphs in SNAP data set [**?**], and λ for one pass. $\lambda_{r,2} = 0.5, \lambda_{r,8} = 0.87$

Data Set	N	nnz	sym	power-law?	$\lambda, k = 2$	$\lambda, k = 8$
amazon0302	262111	1234877	no	yes	0.202	0.37
amazon0312	400727	3200440	no	yes	0.212	0.38
amazon0505	410236	3356824	no	yes	0.206	0.37
amazon0601	403394	3387388	no	yes	0.203	0.36
as-735	7716	26467	yes	yes	0.207	0.44
as-Skitter	1696415	22190596	yes	yes	0.166	0.324
ca-AstroPh	18772	396160	yes	yes	0.232	0.413
ca-CondMat	23133	186936	yes	yes	0.214	0.398
ca-GrQc	5242	28980	yes	yes	0.128	0.241
ca-HepPh	12008	237010	yes	yes	0.082	0.192
ca-HepTh	9877	51971	yes	yes	0.202	0.378
cit-HepPh	34546	421578	no	yes	0.343	0.646
cit-HepTh	27770	352807	no	yes	0.360	0.605
cit-Patents	3774768	16518948	no	yes	0.402	0.726
email-Enron	36692	367662	yes	yes	0.132	0.407
email-EuAll	265214	420045	no	yes	0.280	0.538
Oregon-1	11492	46818	yes	yes	0.224	0.406
Oregon-2	11806	65460	no	yes	0.185	0.429
p2p-Gnutella04	10879	39994	no	yes	0.415	0.747
roadNet-CA	1971281	5533214	yes	no	0.186	0.36
roadNet-PA	1090920	3083796	yes	no	0.188	0.36
roadNet-TX	1393383	3843320	yes	no	0.185	0.358
soc-Epinions1	75888	508837	no	yes	0.173	0.34
soc-LiveJournal1	4847571	68993773	no	yes	0.234	0.463
soc-sign-epinions	131828	841372	no	yes	0.173	0.34
soc-Slashdot0811	77360	905468	no	yes	0.179	0.417
soc-Slashdot0902	82168	948464	no	yes	0.236	0.382
web-BerkStan	685230	7600595	no	yes	0.187	0.301
web-Google	916428	5105039	no	yes	0.189	0.336
web-NotreDame	325729	1497134	no	yes	0.204	0.345
web-Stanford	281903	2312497	no	yes	0.154	0.34
wiki-Talk	2394385	5021410	no	yes	0.411	0.752
wiki-Vote	8297	103689	no	yes	0.409	0.688

The main communication step of the algorithm is the scatter operation. A vector of size n maps each node to a partition, and over the course of the partitioning algorithm, it must be scattered to all nodes. Only $O(n)$ memory must be scattered overall.

Re-streaming also incurs this communication step. Finally, the post-processing step redistributes the graph in what is essentially an all-to-all personalized communication.

As far as I can tell this is the first distributed implementation of streaming partitioning, and if released would be the first publicly available streaming partitioner. Since it's written in C/MPI it would likely be the fastest.

TODO: Benchmark this partitioner within several BFS-intensive kernels on a NERSC machine

2.2.3   Re-streaming Partitioning

An initial idea that Robert Pienta and I came up with was to *re-stream* the partitioning: first, we do a single pass over the nodes, and then we do a second pass, re-assigning nodes to partitions that they have more neighbors to. We could do multiple passes of a streaming partitioner (using the same or different heuristics) to further improve the partitioning, all in a fraction of the time that an offline partitioner would take to terminate. This turned out to be remarkably convergent and effective: after 4 passes, we had improved partition quality from the baseline (1 pass) by a factor of 2-5 (see Fig. **??**). Unfortunately, I discovered a paper that implemented this, after we had already done our own implementation [**?**].


Figure 2: Fraction Improvement in λ after 5 WDG passes, $k = 2$.

Is there an algorithm that traverses a graph multiple times, where it would be worthwhile to 'repartition' and distribute visited vertices in real-time?

2.2.4 Sampling Streaming Partitioning

I implemented a sampling streaming partitioning scheme that randomly samples some percentage of the edges adjacent to any given node, instead of computing the greedy heuristic with *all* neighbors. While the quality of the partitions appeared to be fairly stable, there are two issues that made me decide not to do any further experimentation:

1. The speed benefit was not very large (only 1.2-1.3 times faster), due to unstrided accesses.

2. Since we're streaming the entire graph, we're already going over all of the nnz's, so there's not really any loss in considering all edges.

I guess it's not unthinkable that to find some scenario where this would be worthwhile.

2.3 Thoughts

Ideally we could release the distributed-memory streaming partitioner as a light-weight, open-source package that people could play with / extend, and use in their computations (in place of a heavier partitioner such as METIS).

There are a couple of things left to do:

1. Actually implement support for streaming rather than reading in nodes all at once (this does not change algorithm).

2. Decide if the given input formats (MM / MAT / HB) are acceptable, as well as including the Sparse Matrix Converter as a library.

3. Consider support for parallel file systems.

I think there would be a massive appeal in showing that streaming graph partitioning has a 'killer application,' i.e could be used to accelerate a large computation in industry.

3   Optimizing Distributed Scale-Free Graph Traversals

3.1   Overview

Graph partitioning is already widely used to improve locality of computations, and a wide variety of methods exist [**?**].

Goal: a communication pattern that optimizes distributed graph traversals.

4 Formalizing Characteristics of Graph Traversals

4.1 Overview

4.2 Thoughts

4.3 Results

5 Other Open Questions

Approximate graph algorithms (sampling) – power-reduction

Compact representations of graphs: entropy?

6 TODO

1. Benchmark distributed streaming partitioner within several BFS-intensive kernels on a NERSC machine.

2.